

# Game Playing

# Today's class

- Game playing
  - State of the art and resources
  - Framework
- Game trees
  - Minimax
  - Alpha-beta pruning
  - Adding randomness

# Why study games?

- Clear criteria for success
- Offer an opportunity to study problems involving {hostile, adversarial, competing} agents.
- Historical reasons
- Fun
- Interesting, hard problems which require minimal “initial structure”
- Games often define very large search spaces
  - chess  $35^{100}$  nodes in search tree,  $10^{40}$  legal states

# State of the art

- How good are computer game players?
  - **Chess:**
    - Deep Blue beat Gary Kasparov in 1997
    - Garry Kasparov vs. Deep Junior (Feb 2003): tie!
    - Kasparov vs. X3D Fritz (November 2003): tie!  
<http://www.thechessdrum.net/tournaments/Kasparov-X3DFritz/index.html>
    - Deep Fritz beat world champion Vladimir Kramnik (2006)
  - **Checkers:** Chinook (an AI program with a *very large* endgame database) is the world champion and can provably never be beaten. Retired in 1995
  - **Go:** Computer players have finally reached tournament-level play
  - **Bridge:** “Expert-level” computer players exist (but no world champions yet!)
- Good places to learn more:
  - <http://www.cs.ualberta.ca/~games/>
  - <http://www.cs.unimass.nl/icga>

# Typical case

- 2-person game
- Players alternate moves
- **Zero-sum**: one player's loss is the other's gain
- **Perfect information**: both players have access to complete information about the state of the game. No information is hidden from either player.
- No chance (e.g., using dice) involved
- Examples: Tic-Tac-Toe, Checkers, Chess, Go, Nim, Othello
- Not: Bridge, Solitaire, Backgammon, ...

# How to play a game

- A way to play such a game is to:
  - Consider all the legal moves you can make
  - Compute the new position resulting from each move
  - Evaluate each resulting position and determine which is best
  - Make that move
  - Wait for your opponent to move and repeat
- Key problems are:
  - Representing the “board”
  - Generating all legal next boards
  - Evaluating a position

# Evaluation function

- **Evaluation function** or **static evaluator** is used to **evaluate** the “goodness” of a game position.
  - Contrast with **heuristic search** where the evaluation function was a non-**negative estimate** of the cost from the start node to a goal and passing through the given node
- The zero-sum assumption allows us to use a single evaluation function to describe the goodness of a board with respect to both players.
  - $f(n) \gg 0$ : position  $n$  good for me and bad for you
  - $f(n) \ll 0$ : position  $n$  bad for me and good for you
  - $f(n)$  near  $0$ : position  $n$  is a **neutral position**
  - $f(n) = +\text{infinity}$ : win for me
  - $f(n) = -\text{infinity}$ : win for you


# Evaluation function examples

- Example of an evaluation function for Tic-Tac-Toe:  
$$f(n) = [\text{\# of 3-lengths open for me}] - [\text{\# of 3-lengths open for you}]$$
  
where a 3-length is a complete row, column, or diagonal
- Alan Turing's function for chess
  - $f(n) = w(n)/b(n)$  where  $w(n)$  = sum of the point value of white's pieces and  $b(n)$  = sum of black's
- Most evaluation functions are specified as a weighted sum of position features:  
$$f(n) = w_1 * \text{feat}_1(n) + w_2 * \text{feat}_2(n) + \dots + w_n * \text{feat}_k(n)$$
- Example features for chess are piece count, piece placement, squares controlled, etc.
- Deep Blue had over 8000 features in its evaluation function




# Game trees

- 
- MAX (X)
- MIN (O)
- MAX (X)
- MIN (O)
- TERMINAL
- Utility
- 1 0 +1



# MinMax - Overview

- Search tree
  - *Squares* represent decision states (ie- after a move)
  - *Branches* are decisions (ie- the move)
  - Start at root
  - Nodes at end are leaf nodes
  - Ex: Tic-Tac-Toe (symmetrical positions removed)




- Unlike binary trees can have any number of children
  - Depends on the game situation
- Levels usually called *plies* (a *ply* is one level)
  - Each ply is where "turn" switches to other player
- Players called *Min* and *Max* (next)

# Minimax procedure

- Create **start node as a MAX node** with current board configuration
- **Expand nodes down to some depth** (a.k.a. **ply**) of lookahead in the game
- **Apply the evaluation function at each of the leaf nodes**
- “Back up” values for each of the non-leaf nodes until a value is computed for the root node
  - At **MIN nodes**, the backed-up value is **the minimum of the values** associated with its children.
  - At **MAX nodes**, the backed-up value is the **maximum of the values** associated with its children.
- Pick the operator associated with the child node whose backed-up value determined the value at the root


# Minimax

R




# Minimax

R





# Minimax

**R**



# Minimax Algorithm








# Minimax


Picking my best move against your best move



$minimax(s) =$


$$\begin{cases} \text{utility}(s) & \text{if } \text{terminal}(s) \\ \max_{a \in \text{action}(s)} minimax(\text{result}(s, a)) & \text{if } \text{player}(s) = \text{MAX} \\ \min_{a \in \text{action}(s)} minimax(\text{result}(s, a)) & \text{if } \text{player}(s) = \text{MIN} \end{cases}$$

# Partial Game Tree for Tic-Tac-Toe



- $f(n) = +1$  if the position is a win for X.
- $f(n) = -1$  if the position is a win for O.
- $f(n) = 0$  if the position is a draw.

# Minimax Tree




## Minimax Discussion

- ▶ Complete depth first exploration
- ▶ Depth  $m$  with  $b$  legal moves.  $O(b^m)$
- ▶ Space complexity (memory)  $O(bm)$
- ▶ Chess:  $m \approx 35$ ; on average:  $50 \leq b \leq 100$
- ▶ Impractical for most games, but basis of other algs.

# Alpha-beta pruning

- We can improve on the performance of the minimax algorithm through **alpha-beta pruning**
- Basic idea: *“If you have an idea that is surely bad, don't take the time to see how truly awful it is.”* -- Pat Winston




- We don't need to compute the value at this node.
- No matter what it is, it can't affect the value of the root node.

# Alpha-beta pruning


- Traverse the search tree in depth-first order
- At each **MAX** node  $n$ , **alpha(n)** = maximum value found so far
- At each **MIN** node  $n$ , **beta(n)** = minimum value found so far
  - Note: The alpha values start at -infinity and only increase, while beta values start at +infinity and only decrease.
- **Beta cutoff:** Given a MAX node  $n$ , cut off the search below  $n$  (i.e., don't generate or examine any more of  $n$ 's children) if  $\alpha(n) \geq \beta(i)$  for some MIN node ancestor  $i$  of  $n$ .
- **Alpha cutoff:** stop searching below MIN node  $n$  if  $\beta(n) \leq \alpha(i)$  for some MAX node ancestor  $i$  of  $n$ .

# Alpha-beta example



## Alpha-Beta pruning

### Intuition



Do we need to expand all nodes?

$$\begin{aligned} \text{minimax}(\text{root}) &= \max(\min(\underline{3, 12, 8}), \min(2, x, y), \min(14, 5, 2)) \\ &= \max(3, \min(2, x, y), 2) \\ &= \max(3, z, 2) \\ &= 3 \end{aligned}$$

Do we need  $z$ ?



# Alpha-beta algorithm

```
function MAX-VALUE (state,  $\alpha$ ,  $\beta$ )  
    ;;  $\alpha$  = best MAX so far;  $\beta$  = best MIN  
    if TERMINAL-TEST (state) then return UTILITY(state)  
    v :=  $-\infty$   
    for each s in SUCCESSORS (state) do  
        v := MAX (v, MIN-VALUE (s,  $\alpha$ ,  $\beta$ ))  
        if v  $\geq$   $\beta$  then return v  
         $\alpha$  := MAX ( $\alpha$ , v)  
    end  
    return v
```

```
function MIN-VALUE (state,  $\alpha$ ,  $\beta$ )  
    if TERMINAL-TEST (state) then return UTILITY(state)  
    v :=  $\infty$   
    for each s in SUCCESSORS (state) do  
        v := MIN (v, MAX-VALUE (s,  $\alpha$ ,  $\beta$ ))  
        if v  $\leq$   $\alpha$  then return v  
         $\beta$  := MIN ( $\beta$ , v)  
    end  
    return v
```

Two values:

- ▶  $\alpha$  = value of best choice so far for MAX (highest-value)
- ▶  $\beta$  = value of best choice so far for MIN (lowest-value)
- ▶ Each node keeps track of its  $[\alpha, \beta]$  values

## Alpha-Beta Pruning Properties


- ▶ Pruning does not affect final outcome
- ▶ Sorting moves by result improves  $\alpha - \beta$  performance
- ▶ Perfect ordering:  $O(b^{\frac{m}{2}})$
- ▶ An exercise on **metareasoning**


```
def minimax(position, depth, maximizingPlayer):  
    if depth == 0 or game over in position:  
        return static evaluation of position
```


```
    if maximizingPlayer:  
        maxEval = -infinity  
        for child of position:  
            eval = minimax(child, depth - 1, false)  
            maxEval = max(maxEval, eval)  
        return maxEval
```


```
    else: # minimizingPlayer  
        minEval = +infinity  
        for child of position:  
            eval = minimax(child, depth - 1, true)  
            minEval = min(minEval, eval)  
        return minEval
```

```
return minimax(startPosition, 3, true)
```









```

k(position, depth, alpha, beta, maximizingPlayer)
    or game over in position
    static evaluation of position

```

```

Player
    -infinity
    child of position
    minimax(child, depth - 1, alpha, beta, false)
    = max(maxEval, eval)
    max(alpha, eval)
    <= alpha

```

Eval

```

+infinity
    child of position
    minimax(child, depth - 1, alpha, beta, true)
    = min(minEval, eval)
    min(beta, eval)
    <= alpha


```

Eval

```

11
return Position, 3, -∞, +∞, true)

```





# Effectiveness of alpha-beta

- Alpha-beta is guaranteed to compute the same value for the root node as computed by minimax, with less or equal computation
- **Worst case:** no pruning, examining  $b^d$  leaf nodes, where each node has  $b$  children and a  $d$ -ply search is performed
- **Best case:** examine only  $(2b)^{d/2}$  leaf nodes.
  - Result is you can search twice as deep as minimax!
- **Best case** is when each player's best move is the first alternative generated
- In Deep Blue, they found empirically that alpha-beta pruning meant that the average branching factor at each node was about 6 instead of about 35!

# Games of chance

- Backgammon is a two-player game with **uncertainty**.
- Players roll dice to determine what moves to make.
- White has just rolled *5 and 6* and has four legal moves:
  - 5-10, 5-11
  - 5-11, 19-24
  - 5-10, 10-16
  - 5-11, 11-16
- Such games are good for exploring decision making in adversarial problems involving skill and luck.

