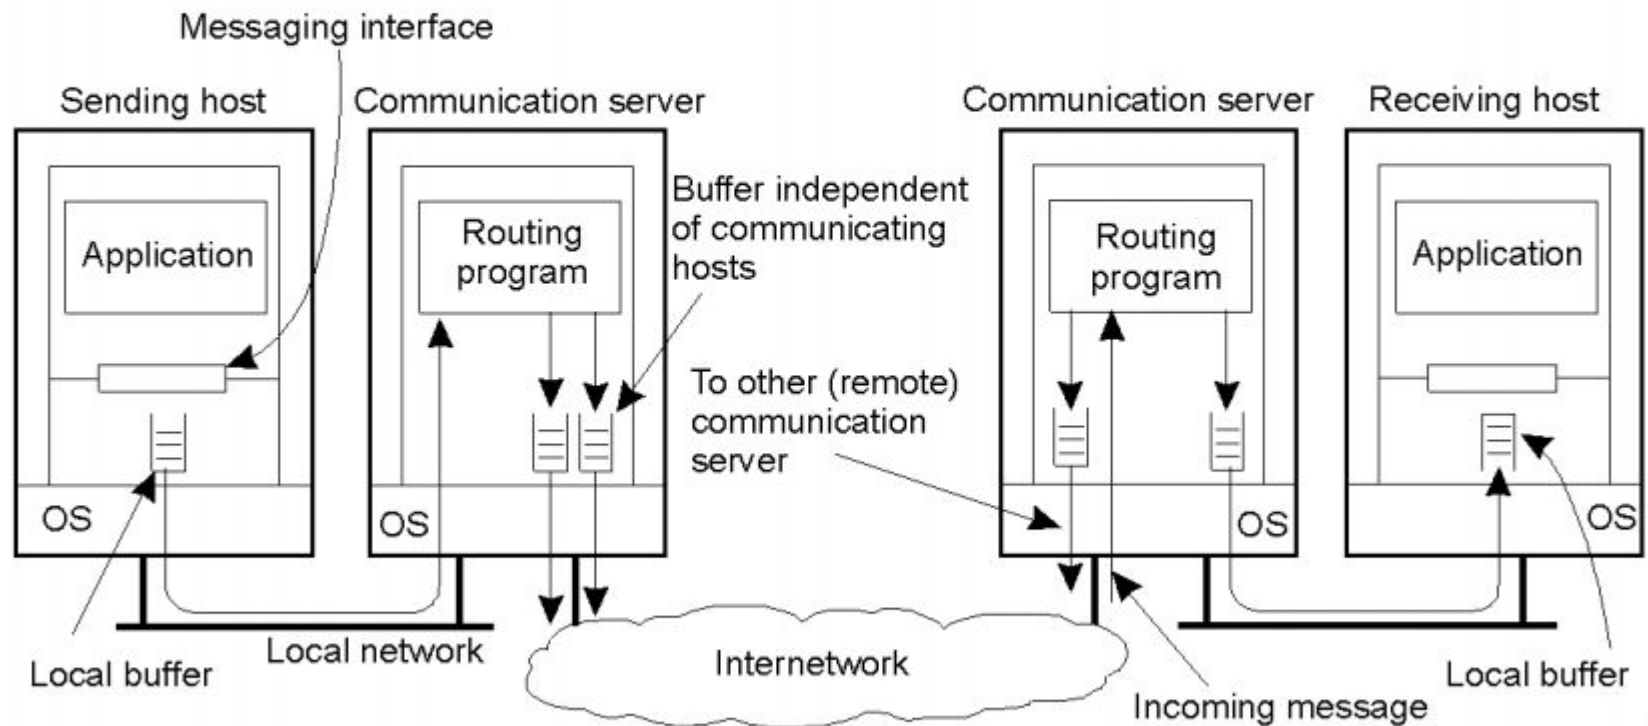


Message-Oriented Communication

Message-Oriented Communication

- Applications are executed on hosts
- The hosts are connected through a network of communication servers
- Each host is connected to some communication server
- **Example: Electronic Mail System**
 - Each host runs an application by which a user can compose, send, receive and read messages.
 - Each host is connected to a mail server
 - Each message is first stored in one of the output buffers of the local mail server.
 - The server removes messages from its buffers and sends them to their destination.
 - The target mail server stores the message in an input buffer for the designated receiver (in the receiver's mailbox).
 - The interface at the receiving host offers a service to the receiver's user agent by which the latter can regularly check for incoming mail.

Persistence and Synchronicity in Communication



Persistence and Synchronicity in Communication

- **Persistent** communication

- a message that has been submitted for transmission is stored by the communication system as long as it takes to deliver it to the receiver.

- **Transient** communication

- a message is stored by the communication system only as long as the sending and receiving application are executing.
 - If a communication server cannot deliver a message to the next communication server or the receiver, the message will be discarded
 - It works like a traditional store-and-forward router

- **Asynchronous** Communication

- A sender continues its execution immediately after it has submitted its message for transmission

- **Synchronous** Communication

- The sender is blocked until its message is stored in a local buffer at the receiving host, or actually delivered to the receiver.

Persistent vs. Transient Communication

Persistent communication: A message is stored at a communication server as long as it takes to deliver it at the receiver.

Transient communication: A message is discarded by a communication server as soon as it cannot be delivered at the next server, or at the receiver.

Synchronous Communication

Some observations: Client/Server computing is generally based on a model of synchronous communication:

- **Client and server have to be active at the time of communication**
- **Client issues request and blocks until it receives reply**
- **Server essentially waits only for incoming requests, and subsequently processes them**

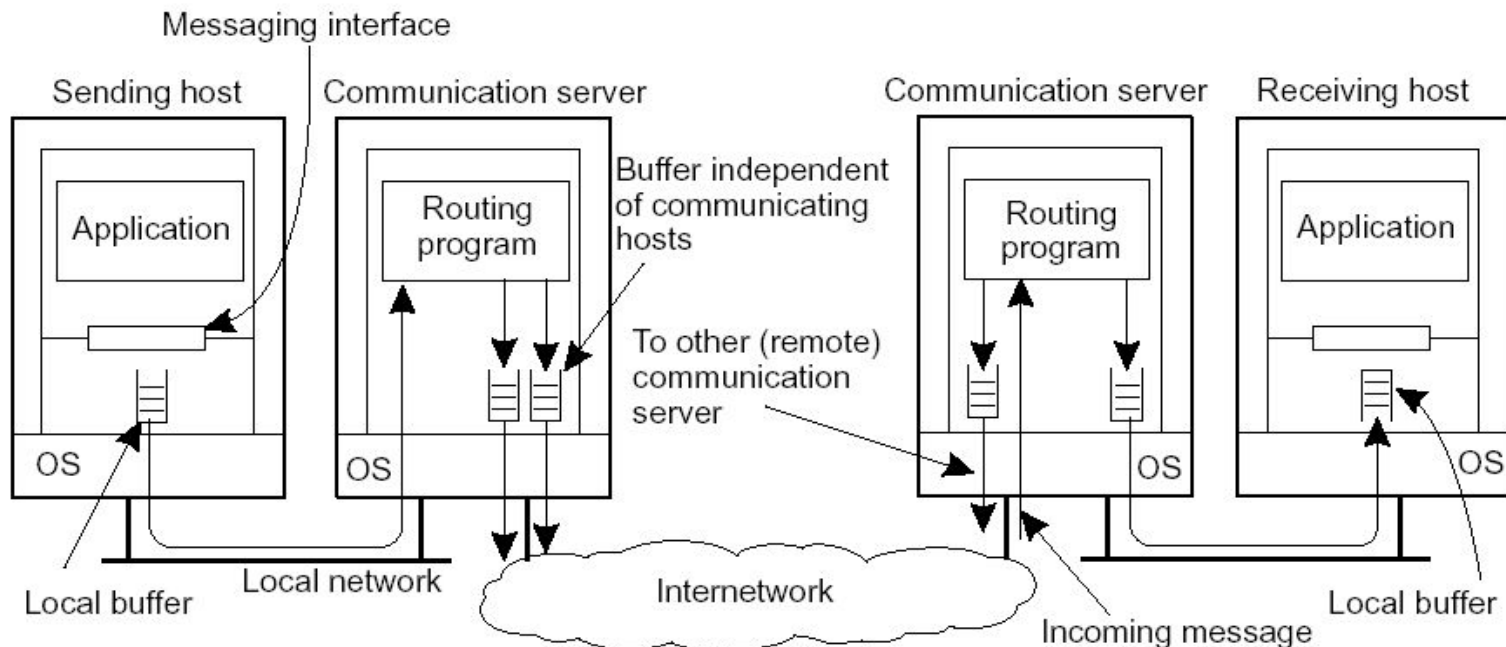
Drawbacks of synchronous communication:

- **Client cannot do any other work while waiting for reply**
- **Failures have to be dealt with immediately (the client is waiting)**
- **In many cases the model is simply not appropriate (mail, news)**

Asynchronous Communication: Messaging

Message-oriented middleware: Aims at high-level asynchronous communication:

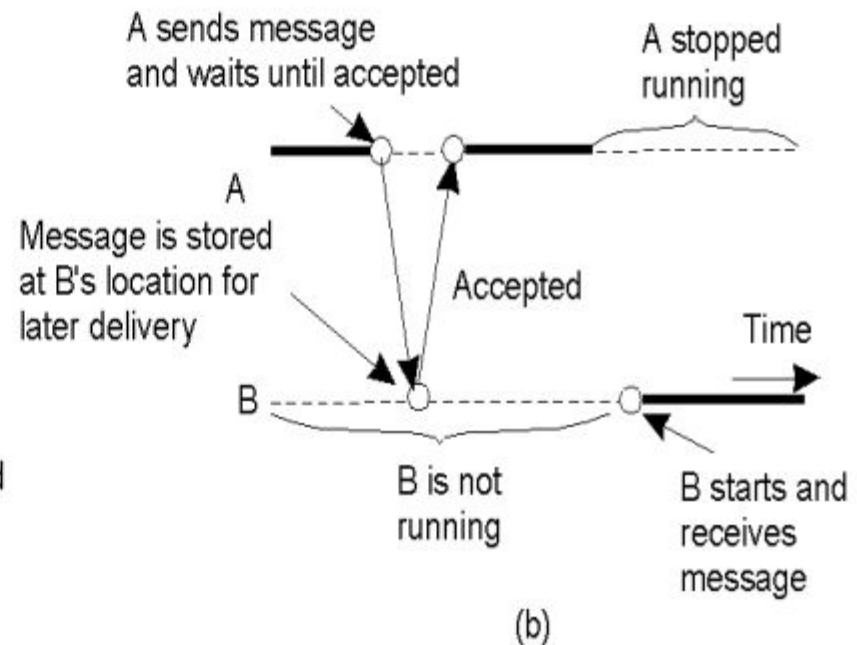
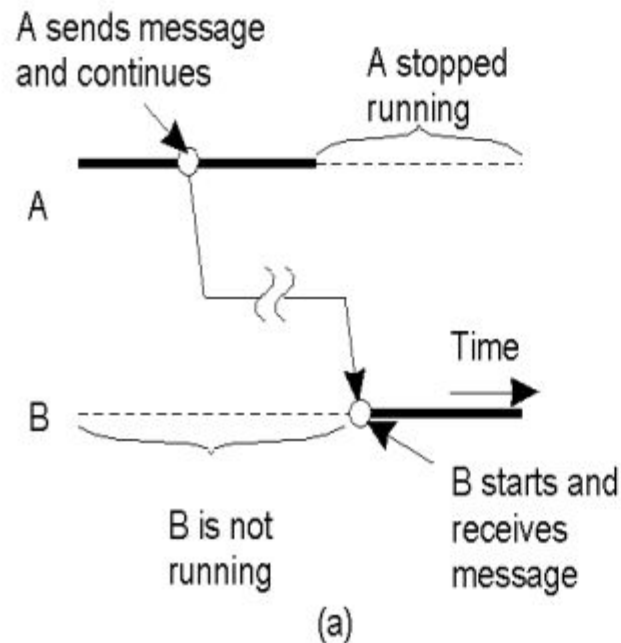
- Processes send each other messages, which are queued
- Sender need not wait for immediate reply, but can do other things
- Middleware often ensures fault tolerance



Persistence and Synchronicity in Communication

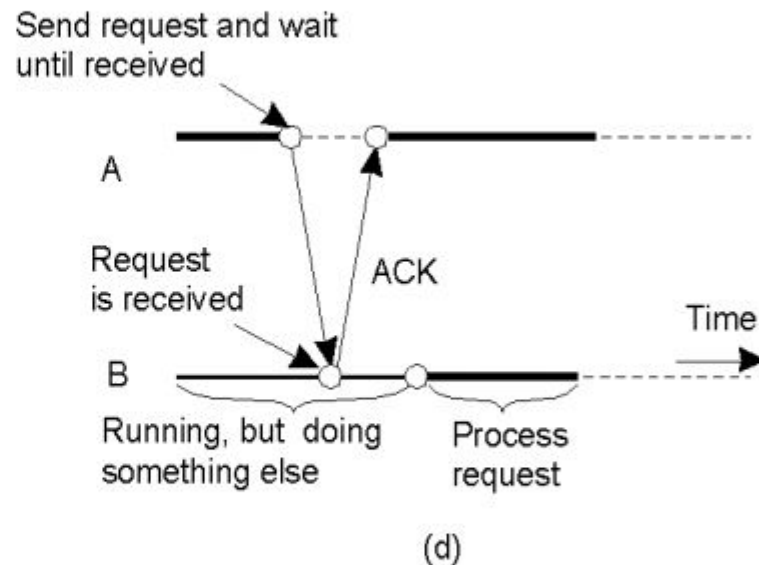

- **Persistent Asynchronous Communication**
 - Each message is either persistently stored in a buffer at the local host or at the first communication server.
 - A e-mail system is an example
- **Persistent Synchronous Communication**
 - Messages can be persistently stored at the receiving host and a sender is blocked until this happens
- **Transient Asynchronous Communication**
 - The message is temporarily stored at a local buffer at the sending host, after which the sender immediately continues
 - UDP is an example
- **Transient Synchronous Communication**
 - The sender is blocked until the message is stored in a local buffer at the receiving host, or
 - until the message is delivered to the receiver for further processing, or
 - until it receives a reply message from the other side (RPCs, RMI)

Persistence and Synchronicity in Communication



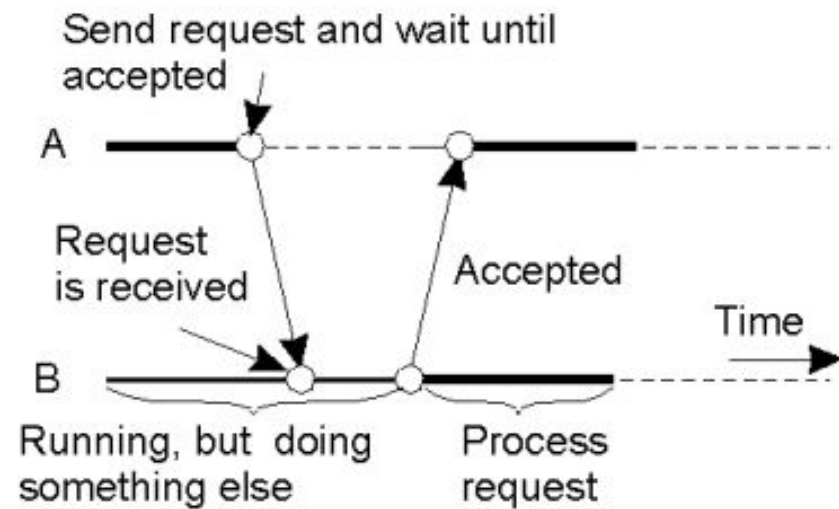
- a) Persistent asynchronous communication
- b) Persistent synchronous communication

Persistence and Synchronicity in Communication

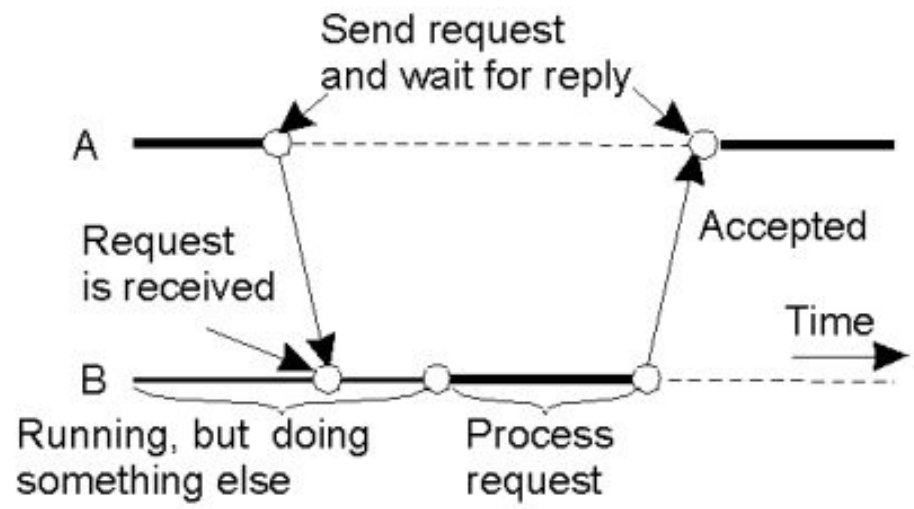


- c) Transient asynchronous communication
- d) Receipt-based transient synchronous communication

Persistence and Synchronicity in Communication



(e)



(f)

- e) Delivery-based transient synchronous communication at message delivery
- f) Response-based transient synchronous communication

Message-Oriented Middleware

Essence: Asynchronous persistent communication through support of middleware-level queues. Queues correspond to buffers at communication servers.

Canonical example: IBM MQSeries

IBM MQSeries (1/3)

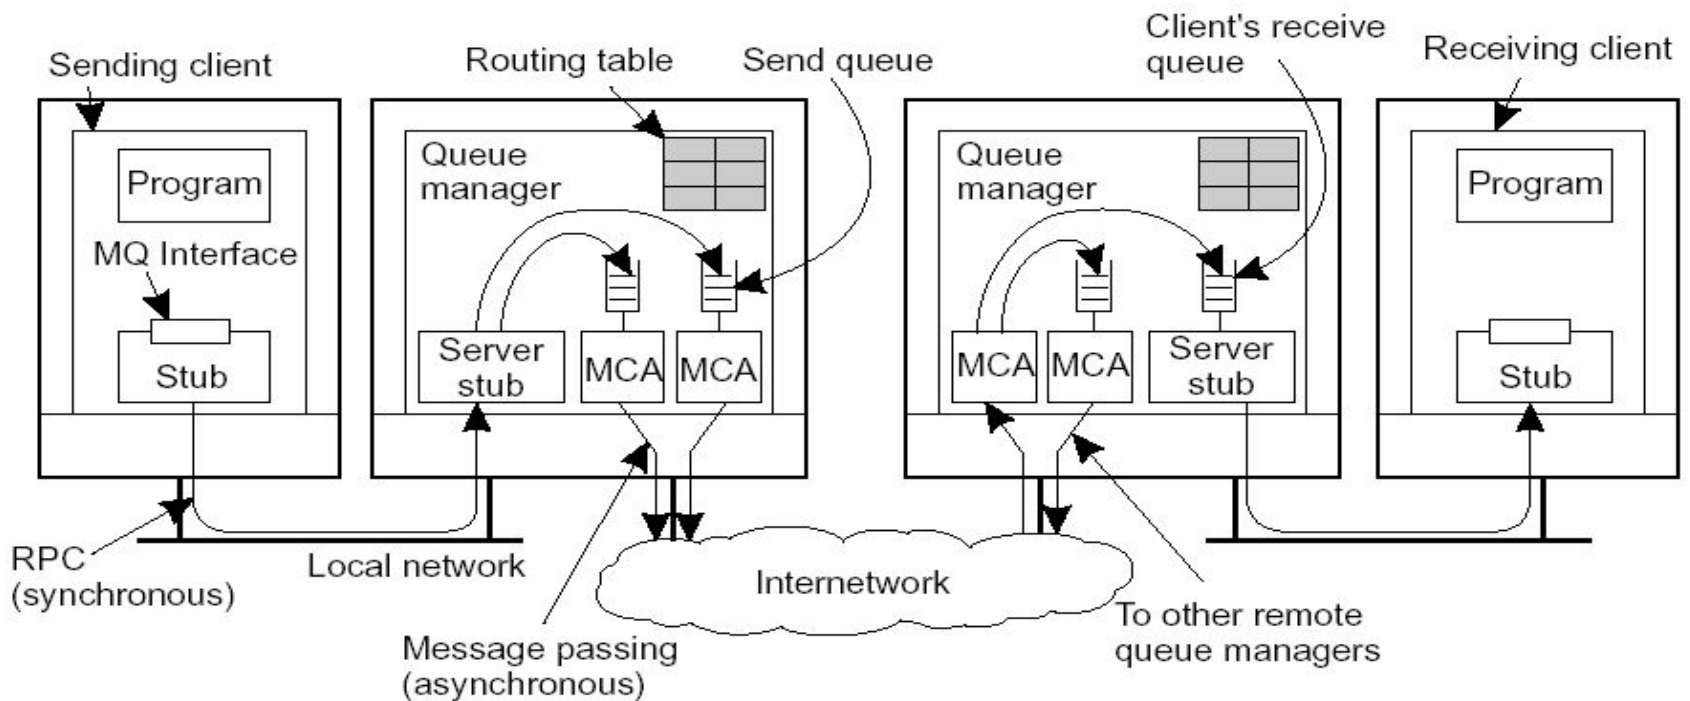
Basic concepts:

- **Application-specific messages are put into, and removed from queues**
- **Queues always reside under the regime of a queue manager**
- **Processes can put messages only in local queues, or through an RPC mechanism**

Message transfer:

- **Messages are transferred between queues**
- **Message transfer between queues at different processes, requires a channel**
- **At each endpoint of channel is a message channel agent (MCA)**
 - **Setting up channels using lower-level network communication facilities (e.g., TCP/IP)**
 - **(Un)wrapping messages from/in transport-level packets**
 - **Sending/receiving packets**


IBM MQSeries (2/3)



- **Channels are inherently unidirectional**
- **MQSeries provides mechanisms to automatically start MCAs when messages arrive, or to have a receiver set up a channel**
- **Any network of queue managers can be created; routes are set up manually (system administration)**

IBM MQSeries (3/3)

Routing: By using logical names, in combination with name resolution to local queues, it is possible to put a message in a remote queue

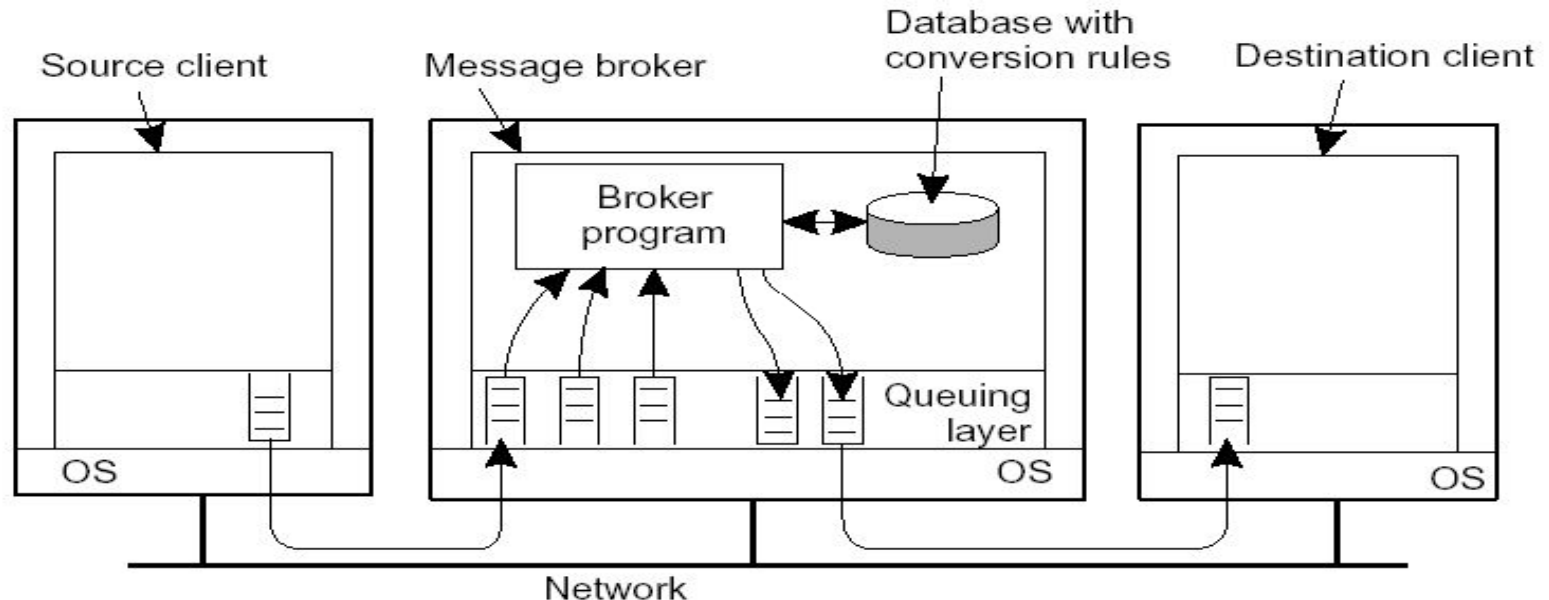


Message Broker

Observation: Message queuing systems assume a *common messaging protocol*: all applications agree on message format (i.e., structure and data representation)

Message broker: Centralized component that takes care of application heterogeneity in a message-queuing system:


- Transforms incoming messages to target format, possibly using intermediate representation
- May provide *subject-based* routing capabilities
- Acts very much like an *application gateway*



What is MPI?

- **A message-passing library specifications:**
 - Extended message-passing model
 - Not a language or compiler specification
 - Not a specific implementation or product
- **For parallel computers, clusters, and heterogeneous networks.**
- **Communication modes: *standard, synchronous, buffered, and ready.***
- **Designed to permit the development of parallel software libraries.**
- **Designed to provide access to advanced parallel hardware for**
 - End users
 - Library writers
 - Tool developers

Group and Context



Group and Context (cont.)

- **Are two important and indivisible concepts of MPI.**
- **Group:** is the set of processes that communicate with one another.
- **Context:** it is somehow similar to the frequency in radio communications.
- **Communicator:** is the central object for communication in MPI. Each communicator is associated with a group and a context.

Communication Modes

- **Based on the type of send:**
 - **Synchronous:** Completes once the acknowledgement is received by the sender.
 - **Buffered send:** completes immediately, unless if an error occurs.
 - **Standard send:** completes once the message has been sent, which may or may not imply that the message has arrived at its destination.
 - **Ready send:** completes immediately, if the receiver is ready for the message it will get it, otherwise the message is dropped silently.

Blocking vs. Non-Blocking

- **Blocking**, means the program will not continue until the communication is completed.
- **Non-Blocking**, means the program will continue, without waiting for the communication to be completed.

Features of MPI

- **General**
 - **Communications combine context and group for message security.**
 - **Thread safety can't be assumed for MPI programs.**

Why to use MPI?

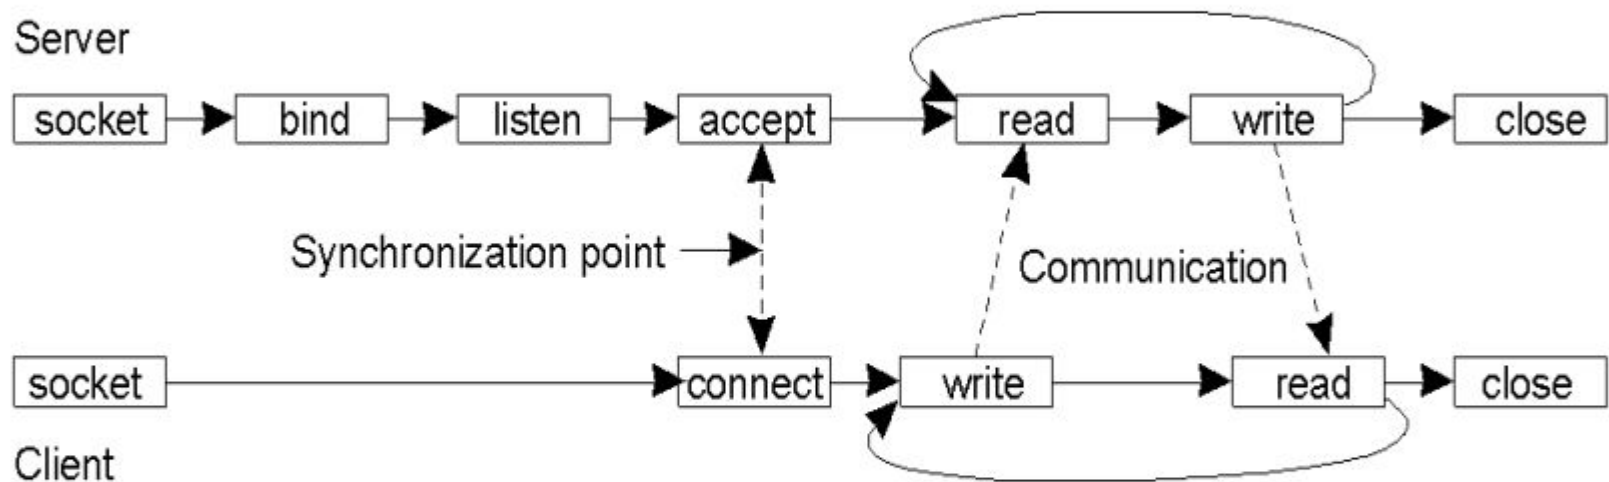
- MPI provides a powerful, efficient, and portable way to express parallel programs.
- MPI was explicitly designed to enable libraries which may eliminate the need for many users to learn (much of) MPI.
- Portable !!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!
- Good way to learn about subtle issues in parallel computing

Berkeley Sockets

| Primitive | Meaning |
|-----------|---|
| Socket | Create a new communication endpoint |
| Bind | Attach a local address to a socket |
| Listen | Announce willingness to accept connections |
| Accept | Block caller until a connection request arrives |
| Connect | Actively attempt to establish a connection |
| Send | Send some data over the connection |
| Receive | Receive some data over the connection |
| Close | Release the connection |

- A socket is a communication endpoint to which an application can write data that are to be sent out over the underlying network, and from which incoming data can be read.
- Socket primitives for TCP/IP.

Berkeley Sockets



- Connection-oriented communication pattern using sockets.